Algorithmic perspectives of kernel method

Chanwoo Lee, August 5, 2020

1 A choice of feature mapping Φ

To derive an algorithm, I choose to use Mapping 1 in the previous note for convenience.

$$\begin{aligned}\Phi\colon \mathbb{R}^{d_1\times d_2} &\to \mathcal{H}_r^{d_1}\times\mathcal{H}_c^{d_2}\\ \boldsymbol{X}&\mapsto (\Phi_r(\boldsymbol{X})\Phi_c(\boldsymbol{X})) \stackrel{\text{def}}{=} (\phi_r([\boldsymbol{X}]_{1:}),\dots,\phi_r([\boldsymbol{X}]_{d_1:}),\phi_c([\boldsymbol{X}]_{:1})\dots,\phi([\boldsymbol{X}]_{:d_2})).\end{aligned}$$

We define decision function,

$$\begin{aligned}f(\boldsymbol{X}) &= \langle \boldsymbol{B}, \Phi(\boldsymbol{X})\rangle, \text{ where } \boldsymbol{B}=(\boldsymbol{B}_r,\boldsymbol{B}_c)\in\mathcal{H}_r^{d_1}\times\mathcal{H}_c^{d_2}\\ &= \langle \boldsymbol{B}_r,\Phi_r(\boldsymbol{X})\rangle + \langle \boldsymbol{B}_c,\Phi_c(\boldsymbol{X})\rangle\\ &= \sum_{k=1}^n \alpha_k^r \sum_{i,j\in[d_2]} w_{ij}^{row} K([\boldsymbol{X}_k]_{i\cdot},[\boldsymbol{X}]_{j\cdot}) + \sum_{k=1}^n \alpha_k^c \sum_{i,j\in[d_2]} w_{ij}^{col} K([\boldsymbol{X}_k]_{\cdot i},[\boldsymbol{X}]_{\cdot j}),\end{aligned} \tag{1}$$

where $\boldsymbol{X}^1,\dots\boldsymbol{X}^n$ are sampled matrix features and $\boldsymbol{W}^{\text{col}},\boldsymbol{W}^{\text{row}}$ are some positive semi definite matrices with low rank. We estimate $\boldsymbol{\alpha}^r=(\alpha_1^r,\dots,\alpha_n^r),\boldsymbol{\alpha}^c=(\alpha_1^c,\dots,\alpha_n^c),\boldsymbol{W}^{\text{col}}$, and $\boldsymbol{W}^{\text{row}}$ from the training data set.

Could you provide the correspondence between (alpha, W) and the outputs in our Algorithm?
We will use W (or P) for visualization/interpretation in data applications —> similar to loading matrices in PCA.

2 Algorithm derivation

We solve an optimization problem

$$\begin{aligned}&\min_{\boldsymbol{B}} \frac{1}{2}\|\boldsymbol{B}\|_F^2 + c\sum_{i=1}^n \xi_i,\\ &\text{subject to } y_i\langle \boldsymbol{B},\Phi(\boldsymbol{X}_i)\rangle \le 1-\xi_i \text{ and } \xi_i\ge 0, i=1,\dots,n.\end{aligned} \tag{2}$$

where $\|\boldsymbol{B}\|_F^2=\|\boldsymbol{B}_r\|_F^2+\|\boldsymbol{B}_c\|_F^2$. From the low rank assumption on $\boldsymbol{B}$ such that

$$\boldsymbol{B}=(\boldsymbol{B}_r,\boldsymbol{B}_c)=\boldsymbol{C}\boldsymbol{P}^T=(\boldsymbol{C}_r,\boldsymbol{C}_c)(\boldsymbol{P}_r,\boldsymbol{P}_c)^T,$$

where $\boldsymbol{C}=(\boldsymbol{C}_r,\boldsymbol{C}_c)\in\mathcal{H}_r^r\times\mathcal{H}_c^r$ and $\boldsymbol{P}=(\boldsymbol{P}_r,\boldsymbol{P}_c)\in\mathbb{R}^{d_1\times r}\times\mathbb{R}^{d_2\times r}$. We assume that $\boldsymbol{P}_r,\boldsymbol{P}_c$ are orthonormal matrices.

1. First we update $\boldsymbol{C}$ holding $\boldsymbol{P}$ fixed. The dual problem of Equation (2) is

$$\begin{aligned}&\min_{\boldsymbol{\alpha}=(\alpha_1,\dots,\alpha_n)} -\sum_{i=1}^n\alpha_i + \frac{1}{2}\sum_{i=1}^n\sum_{j=1}^n \alpha_i\alpha_j y_i y_j \langle \Phi(\boldsymbol{X}_i)\boldsymbol{P}\boldsymbol{P}^T,\Phi(\boldsymbol{X}_j)\boldsymbol{P}\boldsymbol{P}^T\rangle\\ &\text{subject to } \sum_{i=1}^n y_i\alpha_i=0, \text{ and } 0\le\alpha_i\le C, i=1,\dots,n.\end{aligned}$$

Define $\boldsymbol{K}(i,j)\in\mathbb{R}^{d_1\times d_1}\times\mathbb{R}^{d_2\times d_2}$ as

$$\boldsymbol{K}(i,j)=(\boldsymbol{K}_r(i,j),\boldsymbol{K}_c(i,j)) \stackrel{\text{def}}{=} \Phi(\boldsymbol{X}_i)^T\Phi(\boldsymbol{X}_j)$$

$$\text{where } [\boldsymbol{K}_r(i,j)]_{pq} = K_r([\boldsymbol{X}_i]_{p:}, [\boldsymbol{X}_i]_{q:}) \stackrel{\text{def}}{=} \langle \phi_r([\boldsymbol{X}_i]_{p:}), \phi_r([\boldsymbol{X}_j]_{q:}) \rangle,$$
$$[\boldsymbol{K}_c(i,j)]_{pq} = K_c([\boldsymbol{X}_i]_{:p}, [\boldsymbol{X}_i]_{:q}) \stackrel{\text{def}}{=} \langle \phi_c([\boldsymbol{X}_i]_{:p}), \phi_c([\boldsymbol{X}_j]_{:q}) \rangle.$$

Therefore, we can successfully estimate $\boldsymbol{\alpha}$ with quadratic programming based on $\boldsymbol{K}$ without description of feature mapping ϕ_r, ϕ_c. We update $\boldsymbol{C}$ as

$$\boldsymbol{C} = \sum_{i=1}^{n} \alpha_i y_i \Phi(\boldsymbol{X}_i) \boldsymbol{P} \in \mathcal{H}_r^r \times \mathcal{H}_c^r. \tag{3}$$

2. Second, we update $\boldsymbol{P}$ holding $\boldsymbol{C}$ fixed. The dual problem of Equation (2) is

$$\min_{\boldsymbol{\beta}=(\beta_1,\ldots,\beta_n)} -\sum_{i=1}^{n} \beta_i + \frac{1}{2} \sum_{i=1}^{n} \sum_{j=1}^{n} \beta_i \beta_j y_i y_j \langle \boldsymbol{C} \left((\boldsymbol{C}^T \boldsymbol{C})^{-1} \boldsymbol{C}^T \Phi(\boldsymbol{X}_i) \right), \boldsymbol{C} \left((\boldsymbol{C}^T \boldsymbol{C})^{-1} \boldsymbol{C}^T \Phi(\boldsymbol{X}_j) \right) \rangle, \tag{4}$$
$$\text{subject to } \sum_{i=1}^{n} y_i \beta_i = 0, \text{ and } 0 \le \beta_i \le C, i = 1, \ldots, n,$$

Notice $\boldsymbol{C} \left((\boldsymbol{C}^T \boldsymbol{C})^{-1} \boldsymbol{C}^T \Phi(\boldsymbol{X}_i) \right) \in \mathcal{H}^{d_1} \times \mathcal{H}^{d_2}$ is well defined by matrix product: for $\boldsymbol{A}_1 \in \mathcal{H}^r$ and $\boldsymbol{A}_2 \in \mathcal{H}^d$, $\boldsymbol{A}_1^T \boldsymbol{A}_2 = [\![a_{ij}]\!] \in \mathbb{R}^{r \times d}$, where $a_{ij} = \langle [\boldsymbol{A}_1]_i, [\boldsymbol{A}_2]_j \rangle$. We can find an optimizer of (4) without the feature mapping. To show this, notice that by plugging (3) into (4), we have

$$\boldsymbol{C}^T \boldsymbol{C} = \sum_{i=1}^{n} \sum_{j=1}^{n} \alpha_i \alpha_j y_i y_j \boldsymbol{P}^T \boldsymbol{K}(i,j) \boldsymbol{P} \in \mathbb{R}^{r \times r} \times \mathbb{R}^{r \times r}, \tag{5}$$
$$\boldsymbol{C}^T \Phi(\boldsymbol{X}_i) = \sum_{j=1}^{n} \alpha_i y_i \boldsymbol{P}^T \boldsymbol{K}(i,j) \in \mathbb{R}^{r \times d_1} \times \mathbb{R}^{r \times d_2}.$$

(5) makes inner product in (4) expressed in terms of only $\boldsymbol{P}$ and $\{\boldsymbol{K}(i,j) \colon\ i,j \in [n]\}$ by the following equation.

$$\langle \boldsymbol{C} \left((\boldsymbol{C}^T \boldsymbol{C})^{-1} \boldsymbol{C}^T \Phi(\boldsymbol{X}_i) \right), \boldsymbol{C} \left((\boldsymbol{C}^T \boldsymbol{C})^{-1} \boldsymbol{C}^T \Phi(\boldsymbol{X}_j) \right) \rangle = \operatorname{tr} \left(\left(\boldsymbol{C}^T \Phi(\boldsymbol{X}_i) \right)^T \left(\boldsymbol{C}^T \boldsymbol{C} \right)^{-1} \left(\boldsymbol{C}^T \Phi(\boldsymbol{X}_j) \right) \right).$$

Good.

3 Relation with the previous algorithm symmetric trick

Define symmetric feature matrix $\tilde{\boldsymbol{X}} = \begin{pmatrix} 0_{d_1 \times d_2} & \boldsymbol{X} \\ \boldsymbol{X}^t & 0_{d_2 \times d_1} \end{pmatrix} \in \mathbb{R}^{(d_1+d_2) \times (d_1+d_2)}$. Feature mapping 3 is defined as

$$\tilde{\Phi} \colon \mathbb{R}^{d_1 \times d_2} \to \mathcal{H}^{d_1+d_2}$$
$$\boldsymbol{X} \mapsto \left(\phi([\tilde{\boldsymbol{X}}]_{1:}), \ldots, \phi([\tilde{\boldsymbol{X}}]_{d_1+d_2:}) \right)$$

where ϕ is induced by kernel $K \colon \mathbb{R}^{(d_1+d_2) \times (d_1+d_2)} \times \mathbb{R}^{(d_1+d_2) \times (d_1+d_2)} \to \mathbb{R}$. Since all entries of $\Phi_r(\boldsymbol{X})$ are corresponding to $[\tilde{\Phi}(\boldsymbol{X})]_{1:d_1}$ and $\Phi_c(\boldsymbol{X})$ to $[\tilde{\Phi}(\tilde{\boldsymbol{X}})]_{d_1+1:d_1+d_2}$, we have an equivalent representation of (1)

$$f(\boldsymbol{X}) = \langle \boldsymbol{B}, \Phi(\boldsymbol{X}) \rangle$$

$$
\begin{aligned}
&= \langle \boldsymbol{B}_r, \Phi_r(\boldsymbol{X}) \rangle + \langle \boldsymbol{B}_c, \Phi_c(\boldsymbol{X}) \rangle \\
&= \langle \tilde{\boldsymbol{B}}_r, [\tilde{\Phi}(\boldsymbol{X})]_{1:d_1} \rangle + \langle \tilde{\boldsymbol{B}}_c, [\tilde{\Phi}(\boldsymbol{X})]_{d_1+1:d_1+d_2} \rangle, \text{ where } \tilde{\boldsymbol{B}}_r \in \mathcal{H}^{d_1}, \tilde{\boldsymbol{B}}_c \in \mathcal{H}^{d_2} \\
&= \langle \tilde{\boldsymbol{B}}, \tilde{\Phi}(\boldsymbol{X}) \rangle, \text{ where } \tilde{\boldsymbol{B}} = (\tilde{\boldsymbol{B}}_r, \tilde{\boldsymbol{B}}_c) \in \mathcal{H}^{d_1+d_2}.
\end{aligned}
$$

Assume that $\tilde{\boldsymbol{B}} = \tilde{\boldsymbol{C}}\tilde{\boldsymbol{P}}^T$ where $\tilde{\boldsymbol{C}} \in \mathcal{H}^r$, $\tilde{\boldsymbol{P}} = (\tilde{\boldsymbol{P}}_r, \tilde{\boldsymbol{P}}_c) \in \mathbb{R}^{(d_1+d_2)\times r}$ and $\tilde{\boldsymbol{P}}_r \in \mathbb{R}^{d_1 \times r}, \tilde{\boldsymbol{P}}_c \in \mathbb{R}^{d_2 \times r}$. Let Π_r, Π_c are permutation operators such that

$$
\begin{aligned}
\text{Proj}_{\mathcal{H}_r}\left(\Pi_r[\tilde{\Phi}(\boldsymbol{X})]_{1:d_1}\right) &= \Phi_r(\boldsymbol{X}) \\
\text{Proj}_{\mathcal{H}_c}\left(\Pi_c[\tilde{\Phi}(\boldsymbol{X})]_{d_1+1:d_1+d_2}\right) &= \Phi_c(\boldsymbol{X}).
\end{aligned}
$$

Here, we denote $\text{Proj}_{\mathcal{H}_c} \colon \mathcal{H} \to \mathcal{H}_r$ and $\text{Proj}_{\mathcal{H}_r} \colon \mathcal{H} \to \mathcal{H}_c$ as entry-wise projection mappings. Then the following holds

$$
\begin{aligned}
\langle \tilde{\boldsymbol{B}}, \tilde{\Phi}(\boldsymbol{X}) \rangle &= \langle \tilde{\boldsymbol{C}}(\tilde{\boldsymbol{P}}_r, \tilde{\boldsymbol{P}}_c)^T, \tilde{\Phi}(\boldsymbol{X}) \rangle \\
&= \langle \tilde{\boldsymbol{C}}\tilde{\boldsymbol{P}}_r^T, [\tilde{\Phi}(\boldsymbol{X})]_{1:d_1} \rangle + \langle \tilde{\boldsymbol{C}}\tilde{\boldsymbol{P}}_C^T, [\tilde{\Phi}(\boldsymbol{X})]_{d_1+1:d_1+d_2} \rangle \\
&= \langle \Pi_r\tilde{\boldsymbol{C}}\tilde{\boldsymbol{P}}_r^T, \Pi_r[\tilde{\Phi}(\boldsymbol{X})]_{1:d_1} \rangle + \langle \Pi_c\tilde{\boldsymbol{C}}\tilde{\boldsymbol{P}}_c^T, \Pi_c[\tilde{\Phi}(\boldsymbol{X})]_{d_1+1:d_1+d_2} \rangle \\
&= \langle \tilde{\boldsymbol{C}}_r\tilde{\boldsymbol{P}}_r^T, \Phi_r(\boldsymbol{X}) \rangle + \langle \tilde{\boldsymbol{C}}_c\tilde{\boldsymbol{P}}_c^T, \Phi_c(\boldsymbol{X}) \rangle,
\end{aligned}
$$

where $\tilde{\boldsymbol{C}}_r = \text{Proj}_{\mathcal{H}_r}(\Pi_r\tilde{\boldsymbol{C}})$ and $\tilde{\boldsymbol{C}}_c = \text{Proj}_{\mathcal{H}_c}(\Pi_c\tilde{\boldsymbol{C}})$. Therefore, we can conclude that the low rankness of the coefficient on the feature image of $\tilde{\Phi}(\boldsymbol{X})$ implies the same low rankness of the coefficient of the feature image of $\Phi(\boldsymbol{X})$. The other direction is also true.

Good. Updating schemes in both algorithms are also equivalent.